

Midterm 1 Sample Solution

Patrice Belleville

February 9, 2018

NOTE: Throughout the exam a *simple* graph is an undirected, unweighted graph with no multiple edges (i.e., no exact repeats of the same edge) and no self-loops (i.e., no edges from a vertex to itself). Graphs are simple unless stated otherwise, and even where we explicitly contradict one of these, the rest remain true. So, for example, a "directed graph" with no other information specified would be unweighted with no multiple edges and no self-loops.

1 O'd to a Pair of Runtimes [6 marks]

The pairs of functions below represent algorithm runtimes on a pair of lists of length m and n , with $1 < m < n$. For each pair, fill in the circle next to the best choice of:

LEFT: the left function is big- O of the right, i.e., $\text{left} \in O(\text{right})$

RIGHT: the right function is big- O of the left, i.e., $\text{right} \in O(\text{left})$

SAME: the two functions are Θ of each other, i.e., $\text{left} \in \Theta(\text{right})$

INCOMPARABLE: none of the previous relationships holds for all allowed values of n and m .

Do not choose **LEFT** or **RIGHT** if **SAME** is true. The first one is filled in for you.

[2 marks per answer]

Left Function	Right Function	Answer
n	n^2	LEFT
$n \lg m$	$m \lg n$	<input type="radio"/> LEFT <input type="radio"/> RIGHT <input type="radio"/> SAME <input type="radio"/> INCOMPARABLE
$\frac{n^2+m}{m}$	$\frac{m}{\lg m}$	<input type="radio"/> LEFT <input type="radio"/> RIGHT <input type="radio"/> SAME <input type="radio"/> INCOMPARABLE
$n + m \log m$	m^2	<input type="radio"/> LEFT <input type="radio"/> RIGHT <input type="radio"/> SAME <input type="radio"/> INCOMPARABLE

1.1 Solution

1. **RIGHT**: $m \lg n \in O(n \lg m)$ but not vice versa. When we let m grow similarly to n , these are Θ of each other, but when we let n dominate m , $n \lg m$ grows much faster than $m \lg n$.
2. **RIGHT**: We're really comparing $\frac{n^2}{m}$ to $\frac{m}{\lg m}$. $n \in O(\frac{n^2}{m})$, and n dominates $\frac{m}{\lg m}$.
3. **INCOMPARABLE**: When m grows slowly compared to n , the n term can dominate. However, when $m \in \Theta(n)$, then the m^2 term dominates.

2 eXtreme True And/Or False [10 marks]

Each of the following problems presents a scenario and a statement about that scenario. For each one, fill the circle next to the correct one among the three choices:

- The statement is **ALWAYS** true, i.e., true in *every* instance matching the scenario.
- The statement is **SOMETIMES** true, i.e., true in some instance matching the scenario but false in another instance matching the scenario.
- The statement is **NEVER** true, i.e., true in *none* of the instances matching the scenario.

Then, **briefly justify** your answer as follows:

- Justify an **ALWAYS** answer by giving a small instance that fits the scenario for which the statement is **true**, and then briefly sketching the key points in a proof that the statement is **true** for all instances that fit the scenario.
- Justify a **NEVER** answer by giving a small instance that fits the scenario for which the statement is **false**, and then briefly sketching the key points in a proof that the statement is **false** for all instances that fit the scenario.
- Justify a **SOMETIMES** answer by giving **two** small instances that fit the scenario: one for which the statement is **true** and one for which the statement is **false**. (Indicate which one is which!)

Here are the problems: **[5 marks per problem]**

1. **Scenario**: A directed, weakly-connected graph with $n \geq 2$ vertices, m edges and at least two vertices in the same strongly connected component.

Statement: $m \geq n$.

- ☐ **ALWAYS**
☐ **SOMETIMES**
☐ **NEVER**

True instance (always/sometimes) *or* **proof that statement is false in all instances** (never):

False instance (sometimes/never) *or* **proof that statement is true in all instances** (always):

-
2. **Scenario:** An STP (SMP but allowing ties in preference lists) instance in which the preference list of w_1 may contain ties but **no** other preference list contains ties. **Statement:** Gale-Shapley run with **men** proposing on an SMP instance produced by breaking w_1 's ties arbitrarily will produce the same result no matter how the ties are broken.

- ☐ ALWAYS
☐ SOMETIMES
☐ NEVER

True instance (always/sometimes) *or* **proof that statement is false in all instances** (never):

False instance (sometimes/never) *or* **proof that statement is true in all instances** (always):

The rest of the page is intentionally blank.

If you write answers below, **CLEARLY** indicate here what question they belong with **AND** on that problem's page that you have answers here.

2.1 Solution

1. ALWAYS

We can use a two vertex graph like $a \leftrightarrow b$ for our true instance. It has 2 vertices both in the same strongly connected component. It also has two edges.

To prove this, we note that a weakly connected graph must have at least $n - 1$ edges (since when we strip the directions from the edges, it must be connected). However, a weakly connected graph with exactly $n - 1$ edges cannot have two vertices in the same strong component. Again, if we strip the directions, such a graph must be a tree. So, it has no undirected cycles. Every directed cycle is also an undirected cycle when we strip direction; so, with no undirected cycles, we can have no directed cycles and no non-trivial strong components.

2. SOMETIMES

Let's make a true instance. Note that w_1 may have ties but need not actually have any. So, we can just make a standard SMP instance:

```
m1: w1 w2      w1: m1 m2
m2: w2 w1      w2: m2 w1
```

With men proposing, this always ends with $(m_1, w_1), (m_2, w_2)$ as the matching. This is **still** true even if we let the preference list for w_1 be $m_1 = m_2$ instead.

For a false instance, we need w_1 to be "in charge" of the final matching. We can do that by making her the top choice of all men.

```
m1: w1 w2      w1: m1 = m2
m2: w1 w2      w2: m1 w2
```

Now, if we break the ties as m_1, m_2 , then m_1 marries w_1 , but if we break them as m_2, m_1 , then m_2 marries w_1 instead.

3 Oh Oh Oh Oh Oh, Try Everything! [9 marks]

In this problem, we'll investigate what happens when the preconditions of some algorithms are violated.

1. Consider the following `colorize` algorithm from Assignment #2:

```
// G = (V, E) is an undirected graph, represented as an adjacency list
function colorize(V, E)
    n = |V|
    m = |E|
    let Colors be a list of at least n distinct colors
    for i = 0 to n - 1:
        V[i].setvisited(False)

    for i = 0 to n - 1:
        DFS(V[i], (function(v): v.setcolor(Colors[i])))
```

where DFS is defined as:

```
function DFS(v, visitfunction)
    if not v.getvisited():
        visitfunction(v)
        v.setvisited(True)
        for w in neighbours(v):
            DFS(w, visitfunction)
```

When called on an undirected graph, `colorize` colours all vertices in each connected component in the same colour unique to that component. Which of these best completes the following statement about what `colorize` does when called on a **directed** graph? Fill in the circle next to the **best** answer. [2 marks]

When called on a **directed** graph, `colorize` colours all vertices in each...

- ☐ ...connected component in the same colour unique to that component.
- ☐ ...strongly connected component in the same colour unique to that component.
- ☐ ...strongly connected component in the same color but not necessarily unique to that component.
- ☐ None of these.

2. Now, consider the following `intersection2` algorithm from Assignment #2:

```
function intersection2(X, Y)
    Z = { }
    i = j = 0
    m = length(X)
    n = length(Y)
    while i < m and j < n:
        if X[i] < Y[j]:
            i = i + 1
        elsif X[i] > Y[j]:
            j = j + 1
        else:
            add X[i] to Z
            i = i + 1
            j = j + 1
```

With the precondition that X and Y are both sorted in increasing order, this algorithm computes the intersection of X and Y interpreted as sets.

In this problem, we will instead consider the case where $n > 0$, $m = n$, X is sorted, Y is **not necessarily sorted**, and X and Y **contain the same elements** (ignoring order).

- (a) Give exact lower- and upper-bounds on the number of elements in the set `intersection2` returns. Fill in the circle next to **best** answer for each bound. [2 marks]

Lower-bound:

☐ 0

☐ 1

☐ $\lceil \frac{n}{2} \rceil$

☐ n

Upper-bound:

☐ 0

☐ 1

☐ $\lceil \frac{n}{2} \rceil$

☐ n

- (b) Now, give an example with $n = 3$ using your choice of the values 1, 2, 3, and 4 (but no other values) that achieves your **lower-bound**: [3 marks]

$X =$

$Y =$

3. Give a good big-O bound on the runtime of `intersection2` called on a sorted list X and a possibly unsorted list Y of length m and n , respectively. (We have removed the restrictions that $n = m$ and that X and Y contain the same elements.) [2 marks]

Big-O bound:

3.1 Solution

1. **SOLUTION:** When called on a **directed** graph, `colorize` colours all vertices in each strongly connected component in the same color but not necessarily unique to that component.

Why?

A "component" in a directed graph is ill-defined, although a strongly-connected component (or strong component) is well-defined. (Even if we try to give it some plausible definition, it still doesn't match the statement.)

All the nodes in a strongly-connected component **will** be in the same colour. That's because as soon as we hit any one of the nodes for the first time in a DFS, we must hit all of them with the same DFS run (b/c the DFS will hit all reachable nodes, and they're all reachable from each other).

However, that same colour can "bleed" into and out of the strongly-connected component. Consider this little graph: $a \rightarrow b \leftrightarrow c \rightarrow d$. b and c are in the same strong component, but neither a nor d are in that strong component. If we start the first DFS at a , **everything** is the same colour. If we start it at b or c , all but a is the same colour. If we start it at d then next at b or c and finally at a , then each of a , $b + c$, and d gets its own colour.

2. Note that Y need not be sorted, but it **can** be. `intersection2` is correct if both X and Y are sorted. Alternatively, if we "hoist" the biggest element of Y to the front, then we can force `intersection2` to go past almost everything in X before finally finding a match with Y . However, we have no way to force `intersection2` to miss **all** the matches (as long as X and Y contain all the same elements).

That sets us up for our answers:

- (a) Lower-bound on number of returned elements: **1**.
Upper-bound: n
- (b) Here's an instance that generates this: $X = [1, 2, 3]$ and $Y = [3, 1, 2]$. In fact, as long as the basic constraints are followed (whether or not you duplicate elements, by the way), there's a *distinct* largest element, and it goes first in Y , you get only 1 value returned.
3. Whether Y (or X) is in order has no effect on the analysis of `intersection2`. It still makes at least one step of progress in at least one array on each iteration of the loop. So, the maximum number of iterations is $m + n$ and the runtime is in $O(m + n)$ or equivalently $O(\max(m, n))$.

4 Date-A Centre [10 marks]

A cloud service matches processors with jobs to run on them. Each processor is assigned at most one job, each job is assigned to at most one processor, there are n processors p_1, \dots, p_n and n jobs j_1, \dots, j_n , but not all processors can run all jobs. $m(i, k)$ is true if and only if processor p_i can run job j_k .¹

The Cloud Matching Problem (CMP) is to produce a list M such that either $M[k] = X$ (meaning job j_k is not assigned to any processor), or $M[k] = i$ (meaning job j_k is assigned to processor p_i) and $m(i, k)$ is true. Further, in a good solution, each processor is assigned at most one job, and no unused processor can run any unassigned job.

1. Complete the following reduction from CMP to STP (the Stable Marriage Problem with ties).
[6 marks]

Starting with a CMP instance, construct an STP instance as follows:

- Woman w_i is processor p_i .

- Man m_i is

- The preference list for w_i lists

before

- The preference list for m_i lists

before

Then, solve the STP instance to produce a solution S with no strong instabilities. Build a CMP solution by, for each pair (w_i, m_k) in S , setting $M[k]$ to:

¹Perhaps because processor p_i has enough memory, cores, drive space, etc. for job j_k .

2. **Complete the men's and women's preferences** in the following instance to show that the reduction above combined with an STP solver that produces a perfect matching with no strong instabilities (i.e., no instabilities in which m_i and w_j *strictly* prefer each other to their assigned partners) is not enough to **guarantee** that the *largest possible number* of jobs is matched to processors that can run them. [4 marks]

$m(i, k)$ is defined by this table:

	p_1	p_2
j_1	<i>true</i>	<i>true</i>
j_2	<i>true</i>	<i>false</i>

$m_1 : w_1 = w_2 \quad w_1 :$
 $m_2 : \quad \quad \quad w_2 :$

Now **explain the counterexample**:

Without any strong instability in the STP solution, we can assign job j_1 to ☐ p_1 and job ☐ p_2 and job ☐ no processor

j_2 to ☐ p_1 which means that ☐ 0 ☐ 1 ☐ 2 job(s) can be completed. However a better solution ☐ no processor ☐ 2

would be to assign job j_1 to ☐ p_1 and job j_2 to ☐ p_1 so that ☐ 0 ☐ 1 ☐ 2 jobs ☐ p_2 ☐ no processor

can be completed.

4.1 Solution

- Starting with a CMP instance, construct an STP instance as follows:

- Woman w_i is processor p_i .

- Man m_i is

job j_i

- The preference list for w_i lists

jobs that processor i can run

before

jobs that processor i cannot run

- The preference list for m_i lists

processors that can run i

before

processors that cannot run i

Then, solve the STP instance to produce a solution S with no strong instabilities. Build a CMP solution by, for each pair (w_i, m_k) in S , setting $M[k]$ to:

i if $m(i,k)$, and X otherwise

This approach will guarantee us that we'll never have a processor idle and a job unassigned that could have gone together (because such a pair would have constituted a strong instability in STP).

2. We use our reduction to convert this to an STP instance:

$m1: w1 = w2 \quad w1: m1 = m2$
 $m2: w1, w2 \quad w2: m1, m2$

Note that there are two valid solutions here, and neither has strong instabilities. A strong instability requires a pair of unmarried people who mutually *strictly* prefer each other, but there is no pair of people that strictly prefers each other. m_1 and w_1 are simply indifferent (and so neither can pair with anyone to produce a strong instability), and m_2 and w_2 don't like each other.

So, we just have to figure out which solution is worse from a getting jobs done perspective, and that's when we put j_2 onto p_2 .

So:

Without any strong instability in the STP solution, we can assign job j_1 to p₁ and job j_2 to

no processor which means that 1 job(s) can be completed. However a better solution would

be to assign job j_1 to p₂ and job j_2 to p₁ so that 2 jobs can be completed.

5 Not Just for Practice Anymore [5 marks]

Solve these variations on ungraded problems from Assignment #1:

1. Recall that STP is the Stable Marriage problem except that ties are allowed in preference lists. Also recall that a *strong instability* in a perfect matching consists of a woman w and a man m such that w and m both (strictly) prefer each other to their current partners.

Describe a way to create an STP instance of an arbitrary size n that has exactly $(n - 1)!$ solutions without strong instabilities, **briefly** justifying why it has the right number of solutions. **[3 marks]**

2. Imagine that we have a set of n employees. Each one is required to have exactly one mentor chosen from within the same set, but employees can have any number of mentees (zero or more). We can consider any such arrangement to be a directed graph, with each employee as a vertex and an edge from each employee to their mentor.

Which of these is true about any such mentorship graph that follows these rules? Fill in the box next to **every** true statement. **[2 marks]**

- ☐ The graph has at most one cycle.
- ☐ The graph has at least one cycle.
- ☐ Every vertex in the graph is part of a simple cycle.
- ☐ The graph is strongly connected.

5.1 Solution

1. Well, an STP instance with all ties with $n - 1$ people would have exactly $(n - 1)!$ strong stable solutions. If only we could "cut out" one man and woman.

And... we can do just that by saying that m_1 and w_1 strictly prefer each other to everyone else. (Then, pairing them with anyone but each other will generate a strong instability.)

So, we generate an instance with $n \geq 1$ men and women where everyone is entirely indifferent **except** that m_1 ranks w_1 strictly first and w_1 ranks m_1 strictly first.

2. We can easily generate a graph with multiple cycles. For instance, what if e_1 mentors e_2 who mentors e_1 and e_3 mentors e_4 who mentors e_3 .

We can also generate graphs where some vertices are not in simple cycles. How about e_1 mentors e_2 and e_3 . e_2 mentors e_1 . Note that e_3 is not part of any simple cycle.

And neither of these sample graphs was strongly connected.

That leaves having at least one cycle. A pigeonhole argument as in the assignment problem shows that yes, there must be at least one cycle. (Look back at the just-for-practice problems in Assignment #1!)

6 BONUS: From the Cutting Room Floor [2 BONUS marks]

DO ANSWER the new first bonus problem below (but probably don't waste your time on the second one). Otherwise, all directions from the individual exam still apply:

Bonus marks add to your exam and course bonus mark total but are **not** required. **WARNING:** These questions are too hard for their point values. We are free to mark these questions harshly. Finish the rest of the exam before attempting these questions. Do not **taunt** these questions.

1. **ANSWER THIS ONE:** What is one thing you'd like to see on a subsequent assignment in CPSC 320?

To receive 1 bonus point, your answer must be at least as insightful as the response "a picture of a cat", which is a low bar! (Yes, literally writing "a picture of a cat" is worth credit.) We **MAY** allocate one additional course (but not exam) bonus point for especially insightful (or hilarious) responses.

2. **RECALL** from Assignment #2: For its Blue-and-Gold fundraising campaign, UBC has found *anchors*—major donors, each with a *limit* (maximum amount they will donate)—and *causes* to which the anchors will donate. Each cause has a *goal*, the maximum money that will go to the cause. **An anchor donates to at most one cause**, but **one cause may receive donations from many anchors**, and the total amount of the donations of anchors to a cause cannot exceed that cause's goal. The Blue-and-Gold Problem, or BGP, consists in determining how much money each anchor will give to each cause, subject to the constraints stated above. A BGP instance's solution is the maximum dollar amount that can be raised. (Assume limits are distinct, as are goals.)

Consider the following greedy algorithm for this problem:

- (a) Sort the limits in decreasing order.
- (b) For each limit L starting with the largest:

Find the cause with the least money remaining toward its goal (where the amount currently remaining is G) such that $L \leq G$.

If such a cause exists, match L with that cause, and decrease that cause's remaining goal by the minimum of L and G .

Otherwise, find the cause with the largest remaining goal (which may be 0), match L with that cause, and set that cause's remaining goal to 0.

Is this algorithm optimal for this problem?

- ☐ Yes
☐ No

If you answered yes, give a clear, concise, and **complete** proof of your answer (i.e., not just a proof sketch). If you answered no, give a **complete** (but small) counterexample to its correctness (i.e., an instance and also the optimal solution, greedy solution, and clear explanations of each).

6.1 Solution

No solution here! Bonuses are bonuses for a reason :)

Here's some hints, though. No, the algorithm isn't optimal. A way to "trap" the greedy solution is to note that always picking a goal larger than your limit just because it exists is not the best strategy. Sometimes a smaller goal is a better choice.

Oh, and there are pretty small counterexample instances!