

Midterm 1 Sample Solution

NOTE: Throughout the exam a *simple* graph is an undirected, unweighted graph with no multiple edges (i.e., no exact repeats of the same edge) and no self-loops (i.e., no edges from a vertex to itself). Graphs are simple unless stated otherwise, and even where we explicitly contradict one of these, the rest remain true. So, for example, a "directed graph" with no other information specified would be unweighted with no multiple edges and no self-loops.

1 Asymptotically yours

Here are two statements about non-negative functions $f(n)$ and $g(n)$ over the positive integers. For each, indicate whether it is

- **Always** true for every possible pair of non-negative functions $f(n)$ and $g(n)$.
- **Sometimes** true for pairs of non-negative functions $f(n)$ and $g(n)$ (and false for other pairs $f(n)$, $g(n)$).
- **Never** true for any possible pair of non-negative functions $f(n)$ and $g(n)$.

Here are the statements:

1. If $f(n) = cg(n) + o(g(n))$ for some constant c then $f(n) = O(g(n))$.



Always



Sometimes



Never

Additional explanations: $o(g(n))$ is a lower order term, and can be ignored.

2. If $\log f(n) = O(\log g(n))$ then $f(n) = O(g(n))$.



Always



Sometimes



Never

Additional explanations: Note that $\log n^2 = 2 \log n$. So the statement is true for $f(n) = n$ and $g(n) = n^2$, but false for $f(n) = n^2$ and $g(n) = n$.

Now, give examples of nonnegative functions $f(n)$ and $g(n)$ satisfying the following relationships (you should use a different example for each relationship):

3. $f(n) \neq g(n)$ but $f(n) = g(n) + h(n)$ where $h(n) = o(g(n))$.

$$f(n) = \boxed{n + 1} \quad g(n) = \boxed{n}$$

4. $f(n) = O(g(n))$ but $f(n) \neq \Theta(g(n))$.

$$f(n) = \boxed{n} \quad g(n) = \boxed{n^2}$$

5. $f(n) = O(g(n))$ but $2^{f(n)} \neq O(2^{g(n)})$.

$$f(n) = \boxed{2n} \quad g(n) = \boxed{n}$$

2 Knowing your trees

1. Let T be a binary search tree with n keys. In Assignment #1, we considered storing in each node N of the tree the size of the subtree rooted at N , and you described an algorithm that uses this information to count in $O(\log n)$ time the number of keys x of T such that $k_1 < x < k_2$ for some given parameters k_1 and k_2 .

Using the size information stored in each node is only half the job, however. Write appropriate code in each of the blanks in the following implementation of `insert_key` to keep the size information accurate when nodes are inserted in T . Assume that the size field of a node `node` is stored in `node.size`. Note that some of the blanks can (and should) be left blank, and the code does **not** have to maintain balance in the tree.

```
def insert_key(start_node, new_key)
```

```
    if start_node == null:
```

```
        newnode = new node(new_key)
```

```
            newnode.size = 1
```

```
        return newnode
```

```
    if new_key < start_node.key():
```

```
        start_node.size++
```

```
        start_node.left_child = insert_key(start_node.left_child, new_key)
```

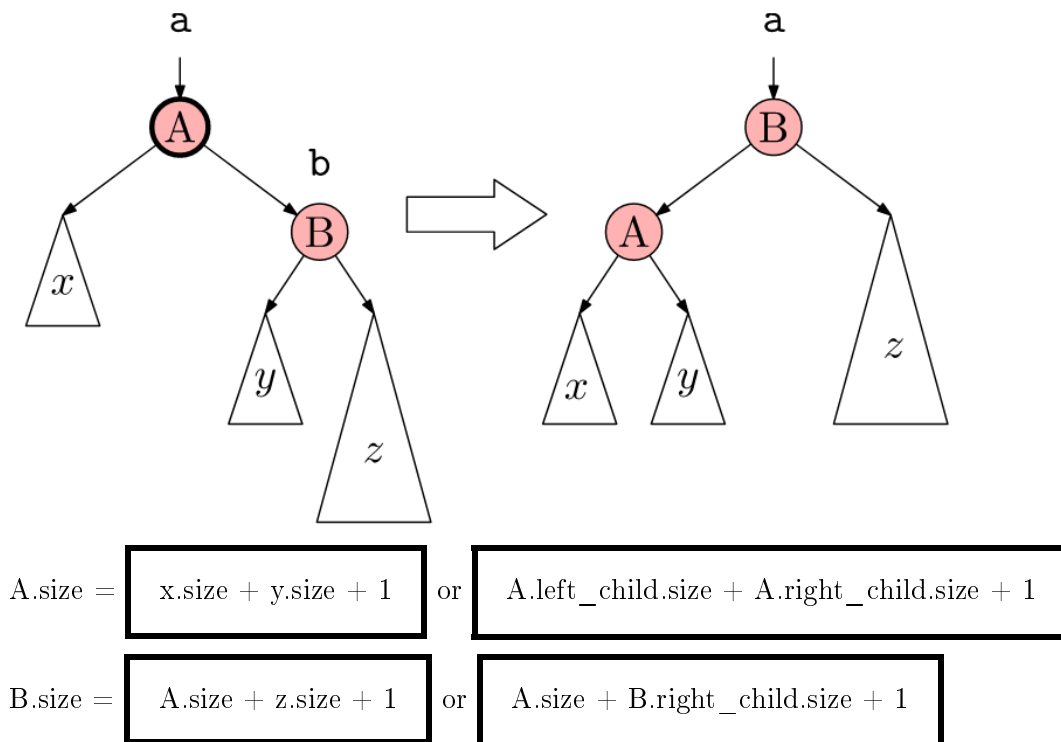
```
    else if new_key > start_node.key():
```

```
        start_node.size++
```

```
        start_node.right_child = insert_key(start_node.right_child, new_key)
```

```
    return start_node
```

2. Now consider a self-balancing AVL tree. Describe how to update the size fields of nodes A and B after a single rotation as illustrated in the figure is performed (we will spare you the other types of rotations). Note that a refers to the parent of the node containing key A , and that x , y and z represent subtrees. Ignore the label b that appears in the picture; it plays no role in this question.



3 More exchange of kidneys

Recall that an instance of the Kidney Exchange Problem (KEP) is (n, M, R) , where n is the number of patients (and also the number of donors), M is an initial perfect matching of patients and donors, and R is a collection of complete rankings $R[p]$ for each of the n patients p . The donor d initially matched with p may not be at the top of p 's ranking.

For an instance I of KEP, the **best-preference graph** G_I has one node per pair (d, p) in the initial matching M . There is a directed edge from (d, p) to (d', p') if d' is the most highly ranked donor on p' 's preference list.

The *Exchange* algorithm of Assignment 2 for matching patients and donors ignores the reality that for large n , it is impractical to get n pairs of patients and donors in the same hospital and do all the exchanges (surgeries) at the same time. In practice, exchanges typically involve just two pairs.

Recall that for a given instance I of KEP, a valid solution S is a perfect matching between patients and donors. S is a *2-exchange* if the following holds for all pairs of patients p and p' : If (p, d) and (p', d') are in the initial matching M and (p, d') is in S , then (p', d) is also in S .

We say that 2-exchange S is *stable* if there do not exist distinct pairs (d, p) and (d', p') such that patient p prefers d' to d **and also** patient p' prefers d to d' .

1. Show an instance of KEP with $n = 3$ for which there is no stable 2-exchange.

Solution : The preference lists are as follows, and the initial matching is $p_1 : d_1$, $p_2 : d_2$, and $p_3 : d_3$.

$p_1 : d_2, d_3, d_1$
 $p_2 : d_3, d_1, d_2$
 $p_3 : d_1, d_2, d_3$

2. Is the number of 2-exchanges upper bounded by a polynomial in n , for all instances I with n patients?
☐ Yes ☒ No

Additional explanations: The number of 2-exchanges is equal to the number of ways we can divide the n patients into pairs. This is $(n-1) \cdot (n-3) \cdots 3 \cdot 1$ which is approximately equal to $\sqrt{n!}$.

3. Are there instances I with n patients for which the number of 2-exchanges exponential in n , for sufficiently large n ?
☒ Yes ☐ No

Additional explanations: Observe that $\sqrt{n!}$ grows faster than any single exponential k^n . So the number of 2-exchanges is (at least) exponential in n .

4. Let *2Exchange* be the following algorithm (a variant of the *Exchange* algorithm from Assignment 2), which arranges exchanges involving two patient-donor pairs, as well as degenerate exchanges involving just one patient-donor pair. The algorithm never arranges exchanges involving three or more pairs.

Algorithm *2Exchange* ($I = (n, M, R)$)

Initialize S to be the empty set

While $n > 0$

 Create the best-preference graph G_I for instance I

 If G_I has a directed cycle, say $v_1, v_2, \dots, v_k = v_1$, for $k = 2$ or $k = 3$ then

 Let v_i be the pair (d_i, p_i) , for $1 \leq i < k$

 Update instance I as follows:

 Remove the pairs (d_i, p_i) from M , $1 \leq i < k$

 Remove d_i from all patient rankings in R , $1 \leq i < k$

 Set n to $n - k + 1$

 Add the pairs (d_{i+1}, p_i) to S for $1 \leq i < k - 1$, and also add the pair (d_1, p_{k-1}) to S

 Else

 Remove all remaining pairs from M and add them to S

 Set n to 0

Endwhile

Return S

Is the output of Algorithm *2Exchange* always the same on any instance, no matter which directed cycle (with $k = 2$ or $k = 3$) is chosen at each iteration of the While loop?

- ☒ Yes ☐ No

Additional explanations: All of the cycles are disjoint, because a vertex has only one outgoing edge. So the order in which we choose them does not matter.

5. Does Algorithm *2Exchange* always produce a 2-stable solution on instances I that have a 2-stable solution?
☐ Yes ☒ No

Additional explanations: Consider four (patient, donor) pairs with the following preferences, and the initial matching $p_1 : d_1, p_2 : d_2, p_3 : d_3$ and $p_4 : d_4$.

$p_1 : d_2, d_3, d_4, d_1$
 $p_2 : d_3, d_4, d_1, d_2$
 $p_3 : d_4, d_1, d_2, d_3$
 $p_4 : d_1, d_2, d_3, d_4$

The best-preference graph for this instance is a four-cycle, which means that algorithm *2Exchange* will do nothing and retain the initial matching. This initial matching is unstable, since p_1 likes d_3 better than d_1 , and p_3 likes d_1 better than d_3 . Thus algorithm *2Exchange* returns a solution that is not 2-stable.

The following solution is 2-stable: $p_1 : d_3, p_2 : d_4, p_3 : d_1, p_4 : d_2$.

4 More on graphs

In this problem, let G be a simple, connected graph with at least two nodes. By the *depth* of a rooted tree, we mean the number of edges of the longest path from the root to a leaf.

For each of the following statements about graph G , indicate whether the statement is:

- Always true for every simple connected graph G with at least two nodes
- Sometimes true for a simple connected graph G , but not for other simple graphs G with at least two nodes
- Never true for any simple graph G with at least two nodes.

Here are the statements:


1. Two different depth first search (dfs) trees rooted at the same node s of G have the same depth.

☐ Always

 ☒ Sometimes

 ☐ Never

Additional explanations: Consider the following graph:



A dfs that starts at B, and visits C first, will give the tree $B - C - D - A - E$ (depth 4). If the dfs visits A first, however, then it will give a tree with depth 3.

2. Two different dfs trees rooted at two *different* nodes s and s' of G have the same depth.

☐ Always

 ☒ Sometimes

 ☐ Never

Additional explanations: Every dfs of the graph shown above that starts at E will give a tree with depth 4; a dfs that starts at A will only give a tree with depth 3.

3. Two different breadth first search (bfs) trees rooted at the same node s of G have the same depth.

☒ Always

 ☐ Sometimes

 ☐ Never

Additional explanations: The depth of a bfs starting at any node N is determined by the node furthest from N in the graph (in terms of the length of the shortest path from N to that node). This does not depend on the order in which the neighbours of each node are visited.

4. Two different bfs trees rooted at two *different* nodes s and s' of G have the same depth.

☐ Always ☒ Sometimes ☐ Never

Additional explanations: Every bfs of the graph shown above that starts at E will give a tree with depth 3; a bfs that starts at A will only give a tree with depth 2.

5 Triathlon reductions

The *Triathlon-stable matching* problem is defined as follows. An instance involves $3n$ athletes: n swimmers, n runners and n cyclists. Each athlete in one group has a preference list for the athletes in each of the other two groups (that is, a runner will have a preference list for the swimmers, and another separate preference list for the cyclists, and so on). You want to find n triples (groups of three), with one swimmer, one runner and one cyclist per triple, to compete in a triathlon. Ideally your solution of n triples should have no instabilities. The first part of this problem asks you to come up with a reasonable notion of instability.

1. Complete the following definition of instability applied to this new situation. With respect to a solution S , an instability consists of two triples (s, r, c) and (s', r', c') of S such that

Solution: An instability is one of the following three situations:

- s likes both r', c' better than r, c respectively **and** r', c' both like s better than s' .
- r likes both s', c' better than s, c respectively **and** s', c' both like r better than r' .
- c likes both r', s' better than r, s respectively **and** r', s' both like c better than c' .

We accepted several other answers (some without any penalty, some with some penalty).

2. Let A be an algorithm that, given the preference lists of all athletes, always finds solution with no instabilities, if the instance has such a solution. Complete the following reduction that uses A to solve the regular stable matching problem:

Given an instance I of Stable Matching we transform it into an instance of Triathlon-stable matching

as follows:

Solution: Here is a reduction that seems reasonable. This was not the only reduction we gave full marks to. Interestingly, we have not yet been able to prove that **any** reduction we came up with works correctly. We'll award bonus points to anyone who manages to come up with a proof of correctness for this or any other reduction.

We perform the following steps:

- we add cyclists c_1, \dots, c_n .
- we rename hospital h_i to swimmer s_i . Its preference list for the r_i 's (renamed from residents to runners) remain the same. Its preference list for the cyclists mirrors its preference list for runners. That is, if r_k occurs in position j on s_i 's preference list, then c_k occurs in position j on s_i 's preference list
- The preference list of r_i for the s_i 's is the same as its preference list for the h_i 's in the instance of stable matching. Its preference list for the cyclists mirrors its preference list for the s_i 's.
- Every cyclist c_i has the same preference list for $r_1 \dots r_n$ as s_i , and the same preference list for $s_1 \dots s_n$ as r_i .